

Full-speed Fuzzing: Reducing Fuzzing Overhead through Coverage-guided Tracing



Stefan Nagy (snagy2@vt.edu), Matthew Hicks (mdhicks2@vt.edu)



Introduction

Of coverage-guided fuzzing's three main components:

(1) *test case generation*, (2) *code coverage tracing*, and (3) *crash triage*, code coverage tracing amounts to **over 90%** of total fuzzer runtime.

Current fuzzers identify coverage-increasing test cases by tracing *all of them*—even when **over 99.99%** do not increase coverage and thus are discarded.

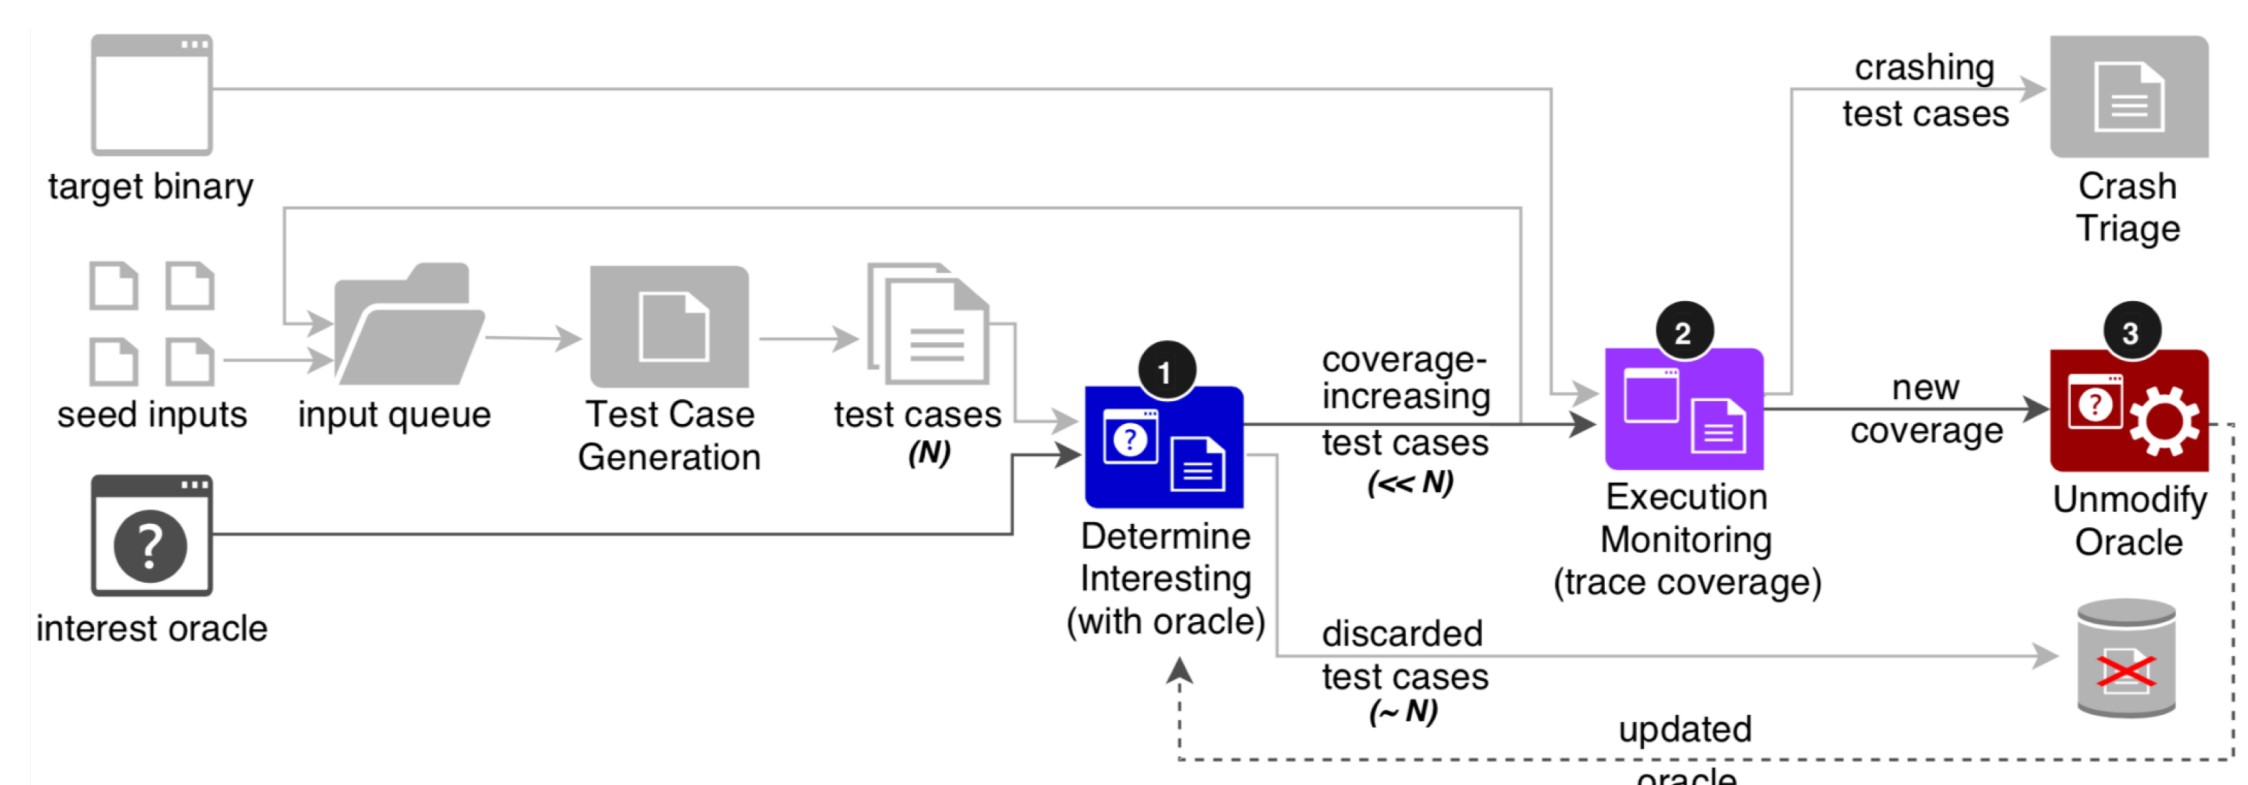
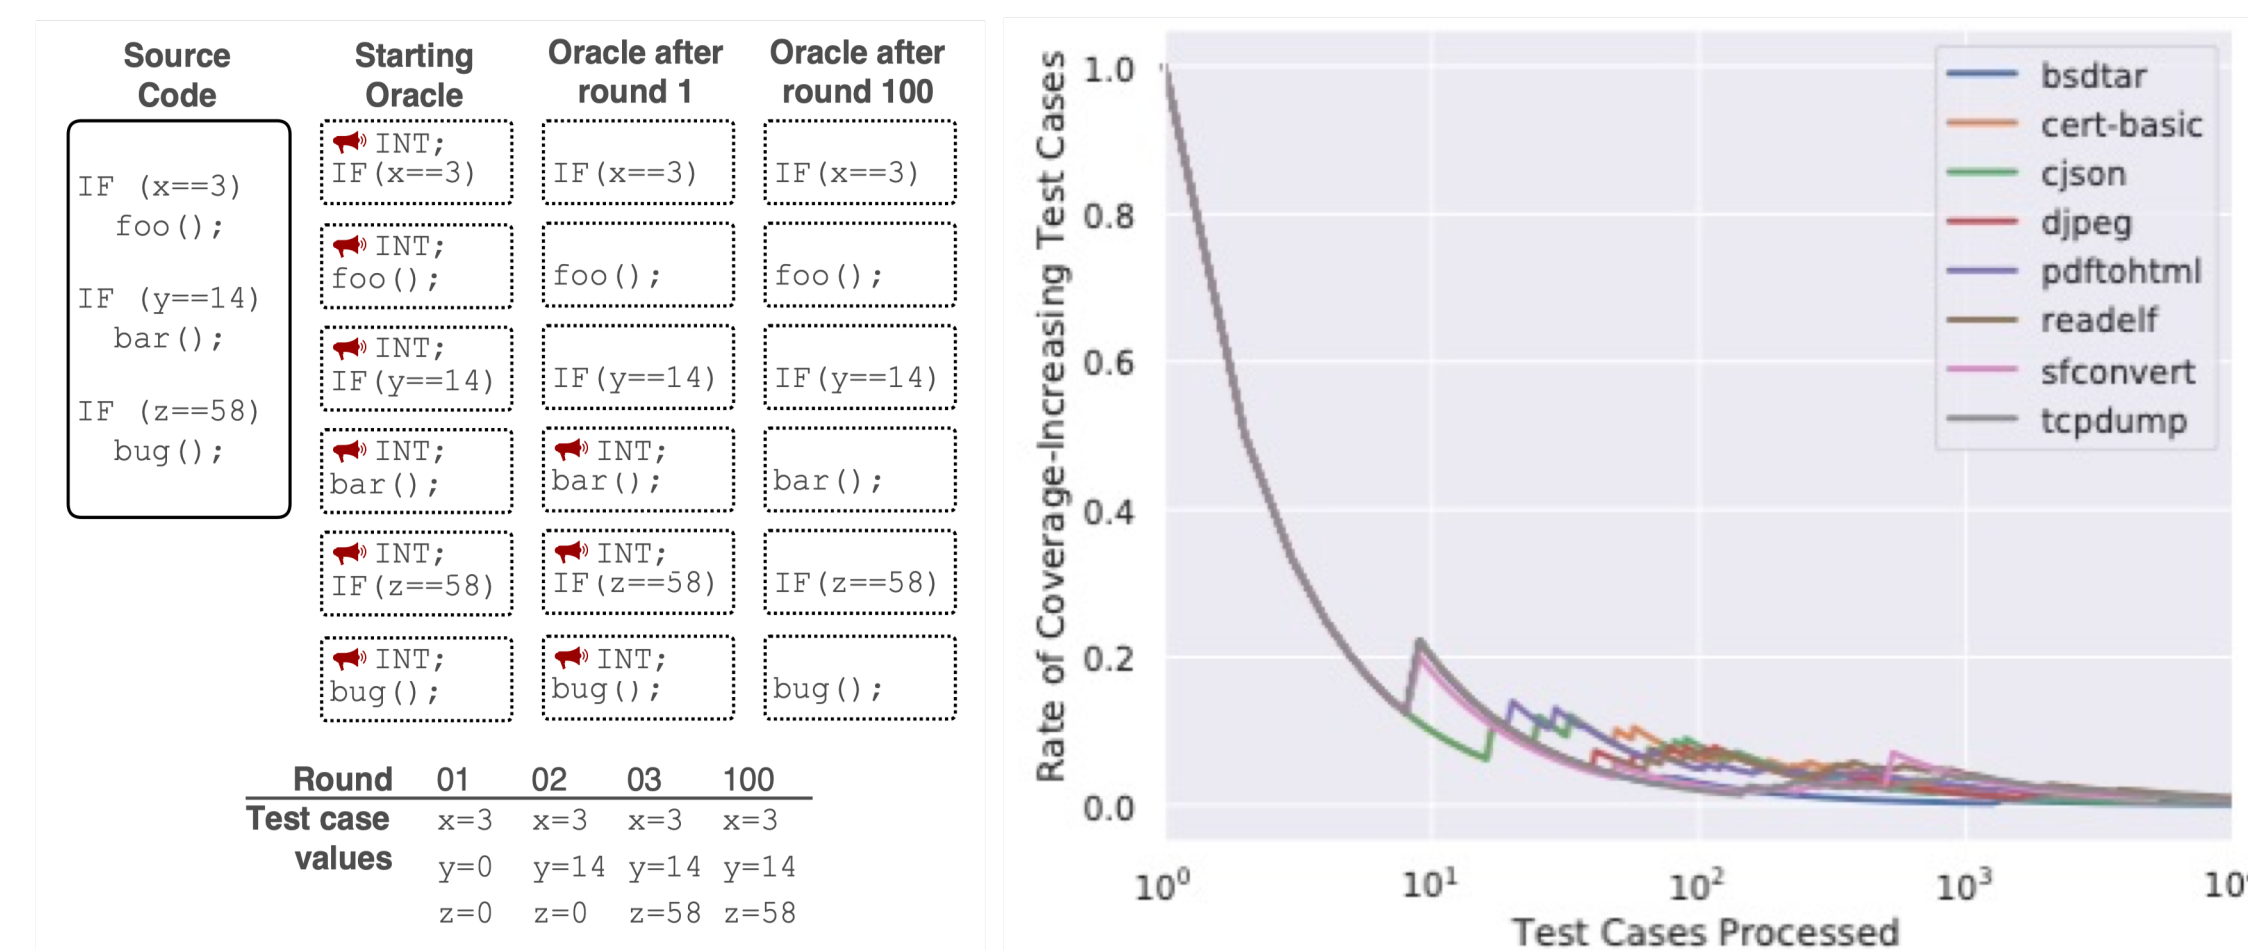
We introduce **Coverage-guided Tracing**—an approach restricting tracing to test cases *guaranteed* to increase coverage while filtering-out the rest at *native speed*.

Coverage-guided Tracing


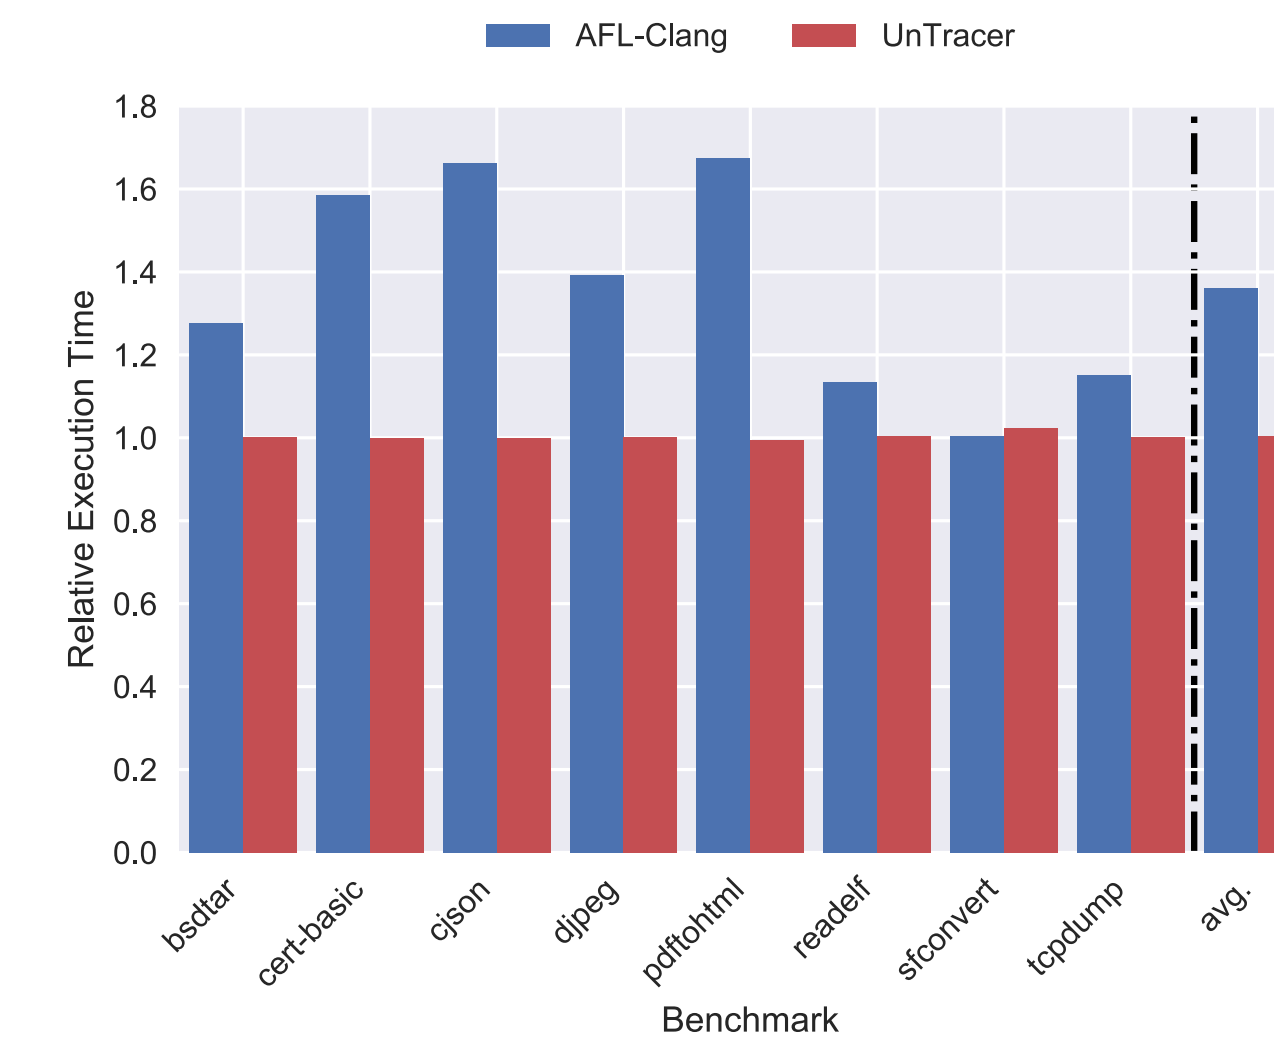

Driving observations:

- (1) a small fraction of generated test cases increase coverage
- (2) coverage-increasing test cases become less frequent over time

Coverage-guided tracing transforms the target binary so that it self-reports when a test case produces new coverage—without tracing. This restricts the expense of tracing to *only* coverage-increasing test cases.



Tracing-only Overhead and Hybrid Fuzzing Evaluations



UnTracer: coverage-guided tracer based on the Dyninst black-box binary static rewriter


Tracing-only overhead evaluation:

- **36%** faster than AFL-Clang.
- **518%** faster than AFL-Dyninst.
- **612%** faster than AFL-QEMU.

Hybrid fuzzing throughput evaluation:

- **79%** more test cases QSYM-Clang.
- **618%** more test cases than QSYM-QEMU.

Discussion



- Orthogonal to other improvements in fuzzing.
- Increased overhead for coverage-increasing test cases ($\ll N$) amortized by near-native speed for the rest ($\sim N$).
- Fully black-box (binary-only) approaches are feasible.
- Edge coverage an open challenge.

Conclusions & Future Work

Coverage-guided tracing leverages the fact that coverage-increasing test cases are the overwhelmingly uncommon case in fuzzing by modifying target binaries so that they self-report when a test case produces new coverage.

We report overhead reductions of as much as 1300% and 70% for black- and white-box tracing, respectively, and 616% and 79% more test case executions than black- and white-box tracing-based hybrid fuzzing.

Our current work focuses on improving the performance and precision of black-box binary fuzzing via static rewriting.

Where to find our software?



www.github.com/FoRTE-Research

- **UnTracer-AFL** : UnTracer integrated with AFL
- **afl-fid** : AFL for fixed-dataset experiments
- **FoRTE-FuzzBench** : our eight fuzzing benchmarks



Scan me



COMPUTER SCIENCE
VIRGINIA TECH