

Bin2Wrong:

A Unified Fuzzing Framework for Uncovering Semantic Errors in Binary-to-C Decompilers

Zao Yang

University of Utah


Stefan Nagy

University of Utah



Decompilers: Critical to Software & Systems

- High-level goal: recover equivalent C code from compiled binary artifacts



Compiled Binaries



Decompilers

```
int func1(int x) {
    int result = 0;
    if (x > 10) {
        result = x * 2;
    } else {
        result = x + 5;
    }


    for (int i = 0; i < 3; i++) {
        result += i;
    }

    return result;
}
```

Recovered C Code

Decompilers: Critical to Software & Systems

- **High-level goal:** recover equivalent C code from compiled binary artifacts
 - Fundamental to downstream tasks that center on **source-unavailable components**



Proprietary Software
Performance Tuning



Obfuscated Binary
Malware Analysis




Commercial Software
Vulnerability Discovery

Decompilers: Critical to Software & Systems

- **High-level goal:** recover equivalent C code from compiled binary artifacts
 - Fundamental to downstream tasks that center on **source-unavailable components**
 - **Success of these downstream tasks** often undermined by **incorrect decompilation**



Buggy Decompilers



Wrong Semantics



Misled Users



Unsuccessful Tasks

Decompilers: Critical to Software & Systems

- **High-level goal:** recover equivalent C code from compiled binary artifacts
 - Fundamental to downstream tasks that center on **source-unavailable components**
 - **Success of these downstream tasks** often undermined by **incorrect decompilation**



Buggy Decompilers



Wrong Semantics



Misled Users



Unsuccessful Tasks

Problem: what factors chiefly influence **decompilation errors?**

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**

Origin(s)	Affected Decompilers
Instructions	Ghidra [1], Radare2 [2], Reko [3], RetDec [4, 5]
Refinement	Angr [6, 7], Binary Ninja [8, 9, 10], RetDec [11]
If / Else	Binary Ninja [12, 13, 14], Ghidra [15], RetDec [16, 17]
Loops	Angr [18, 19, 20], Binary Ninja [21, 22], Reko [23]
Goto	Angr [24, 25], Binary Ninja [26], Reko [27], RetDec [28]
Switches	Binary Ninja [29, 30, 31], Radare2 [32], RetDec [33]
Arguments	Angr [34, 35, 36], Binary Ninja [37, 38, 39, 40]
Variables	Angr [41, 42], Binary Ninja [43], Ghidra [44, 45, 46]
Literals	Angr [47], Binary Ninja [48], Ghidra [49, 50], Reko [51]
Compilers, Opts	Angr [52, 53, 54], Binary Ninja [55], Ghidra [56, 57]
Executable Formats	Angr [58], Binary Ninja [59, 60, 61], Ghidra [62, 63]

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**

```
int v0 =  
divide(nmr, dnr);
```



```
word32 v01 =  
sub_4a(v42, v13);
```

Opaqueness of
Source Semantics

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**

```
int v0 =  
divide(nmr, dnr);
```



```
word32 v01 =  
sub_4a(v42, v13);
```

```
mov eax, edi  
add eax, eax
```



```
mov eax, edi  
shl eax
```



```
mov eax, edi  
shl eax, 0x1
```



Opaqueness of
Source Semantics

Patterns among
Different Compilers

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**

```
int v0 =  
divide(nmr, dnr);
```

```
mov eax, edi  
add eax, eax
```



```
word32 v01 =  
sub_4a(v42, v13);
```

```
mov eax, edi  
shl eax
```



```
mov eax, edi  
shl eax, 0x1
```



Opaqueness of
Source Semantics

Patterns among
Different Compilers

```
for (i=0;i<3;i++)  
a[i]--;
```

```
a[0]--;  
a[1]--;  
a[2]--;
```

Layout-altering
Code Optimizations

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**

```
int v0 =  
divide(nmr, dnr);
```

```
mov eax, edi  
add eax, eax
```



```
word32 v01 =  
sub_4a(v42, v13);
```

```
mov eax, edi  
shl eax
```



```
mov eax, edi  
shl eax, 0x1
```


Opaqueness of
Source Semantics

Patterns among
Different Compilers

```
for (i=0;i<3;i++)  
a[i]--;
```

```
a[0]--;  
a[1]--;  
a[2]--;
```

Layout-altering
Code Optimizations



Differences in
Executable Formats

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**
 - Root causes stem not just from **individual factors**—but **combinations** as well

```
if (var != 0xF)  
if (0xF != 0xF)
```

Incorrect **if** condition

```
if (x < 1.5)  
if (1.5 < x)
```

Mi-swapped **operands**

```
0.00000000  
0.0000000023283
```

Erroneous **float** value

```
void myFunc(int a, float b)  
void myFunc(int a [ ] )
```

Missed **float** argument in PE binaries

```
return tailFunc(input);  
return 0;
```

Dropped **call** in tailcall optimization

Opaqueness of
Source Semantics

Patterns among
Different Compilers

Layout-altering
Code Optimizations

Differences in
Executable Formats

Decompilers: Challenges to Accurate Decompilation

- Manually distilled **64** prior decompiler bugs into **four distinct factors**
 - Root causes stem not just from **individual factors**—but **combinations** as well

```
if (var != 0xF)  
if (0xF != 0xF)
```

Incorrect **if** condition

```
if (x < 1.5)  
if (1.5 < x)
```

Mi-swapped **operands**

```
0.00000000  
0.0000000023283
```

Erroneous **float** value

```
void myFunc(int a, float b)  
void myFunc(int a [ ] )
```

Missed **float** argument in **PE** binaries

```
return tailFunc(input);  
return 0;
```

Dropped **call** in **tailcall** optimization

- Preventing such errors demands **thorough testing along these factors**

Prior Work: Automated Decompiler Testing

- **Problem:** prior approaches only **marginally explore source & compilation**
 - Hardcoded to **limited source constructs**, **specific compilers/optimizations**, and **just ELF**

Approach	Source Construct Diversity			Compilation Configuration Diversity					Decompiler Agnostic
	Expressions	Control Flows	Data	Compilers	Optimizations	ELF	PE	MachO	
DecFuzzer	✓	👎	👎	1	✗	✓	✗	✗	✓
Cornucopia	✓	✓	✓	2	✓	✓	✗	✗	✓
D-Helix	✓	✓	👎	2	👎	✓	✗	✗	✗
DSmith	✓	👎	👎	1	👎	✓	✗	✗	✓

Key: ✓ = fully support, ✗ = no support, 👎 = limited constrained support

Prior Work: Automated Decompiler Testing

- **Problem:** prior approaches only **marginally explore source & compilation**
 - Hardcoded to **limited source constructs**, **specific compilers/optimizations**, and **just ELF**


Approach	Source Construct Diversity			Compilation Configuration Diversity				Decompiler Agnostic
	Expressions	Control Flows	Data	Compilers	Optimizations	ELF	PE	
DecFuzz	✓	✓	✓	✓	✓	✓	✓	✓
Cornucopia	✓	✓	✓	✓	✓	✓	✓	✓
D-Helix	✓	✓	✓	✓	✓	✓	✓	✗
DSmith	✓	👎	👎	1	👎	✓	✗	✓

Motivation: how can fuzzing systematically explore these **individual** binary factors and **their many combinations?**

Key: ✓ = fully support, ✗ = no support, 👎 = limited constrained support


Our Solution: Bin2Wrong

- **Idea:** mutate source code, compilers, optimizations, and format **altogether**
 - **Backbone:** multi-dimensional test case format
 - Mutate **compilation configuration** via byte flips
 - Mutate **source layout** via AST-level mutations
- **Out-of-the-box support for:**
 - **Compilers:** GCC, Clang, ICX, TCC, AppleClang, MSVC
 - **Optimizations:** 5,183 total across all compilers
 - **Formats:** Linux ELF, MacOS MachO, Windows PE
 - **Source:** all C constructs incl. strings/floats/gotos




Our Solution: Bin2Wrong

- Idea: mutate source code, compilers, optimizations, and format **altogether**




Our Solution: Bin2Wrong

- Idea: mutate source code, compilers, optimizations, and format **altogether**




Our Solution: Bin2Wrong

- Idea: mutate source code, compilers, optimizations, and format **altogether**



Our Solution: Bin2Wrong

- Idea: mutate source code, compilers, optimizations, and format **altogether**



Our Solution: Bin2Wrong

- Idea: mutate source code, compilers, optimizations, and format **altogether**

Approach	Source Construct Diversity			Compilation Configuration Diversity					Decompiler Agnostic
	Expressions	Control Flows	Data	Compilers	Optimizations	ELF	PE	MachO	
DecFuzzer	✓	👎	👎	1	✗	✓	✗	✗	✓
Cornucopia	✓	✓	✓	2	✓	✓	✗	✗	✓
D-Helix	✓	✓	👎	2	👎	✓	✗	✗	✗
DSmith	✓	👎	👎	1	👎	✓	✗	✗	✓
Bin2Wrong	✓	✓	✓	6	5,183	✓	✓	✓	✓

By maximizing source *and* compilation diversity, Bin2Wrong enables the most thorough evaluation of decompilation correctness to date.

Evaluation: Overview

■ Fundamental questions:

- Does Bin2Wrong's unified mutation attain **greater binary diversity?**
- Can Bin2Wrong-generated binaries test **more decompiler internals?**
- Will Bin2Wrong's binaries discover **more decompilation bugs?**

■ Competing decompiler fuzzers:

- Corncopia (mutates optimizations)
- DecFuzzer (mutates source code)

Open-source Decompilers	 Angr	 Radare2
	 Relyze	 RetDec
	 Reko	 Rev.Ng
Commercial Decompiler	 Binary Ninja	

Benchmarked Decompilers

Evaluation: Binary Diversity

- Measured **mean binary-to-binary diversity** across 10,000 binaries
 - Similarity scoring calculated via **three** state-of-the-art diffing algorithms

Diffing Algorithms
BinDiff (Zynamics's Graph-Based)
Radiff2-M (Eugene W. Myers' $O(ND)$)
Radiff2-L (Levenshtein's Edit Distance)

Evaluation: Binary Diversity

- Measured **mean binary-to-binary diversity** across 10,000 binaries
 - Similarity scoring calculated via **three** state-of-the-art diffing algorithms

Diffing Algorithms	Bin2Wrong vs. Cornucopia	Bin2Wrong vs. DecFuzzer
BinDiff (Zynamics's Graph-Based)	+9.398 X	+16.189 X
Radiff2-M (Eugene W. Myers' O(ND))	+8.119 X	+15.941 X
Radiff2-L (Levenshtein's Edit Distance)	+6.131 X	+15.089 X

- Takeaways:** Bin2Wrong's unified mutation greatly **increases binary diversity**

Evaluation: Decompiler Code Coverage

- Measured **mean decompiler code coverage** across 24-hour fuzzing trials
 - Computed basic block coverage via the AFL-QEMU-Cov utility

Metrics	Bin2Wrong vs. Cornucopia	Bin2Wrong vs. DecFuzzer
Basic Block Coverage	+16%	+32%

Evaluation: Decompiler Code Coverage

- Measured **mean decompiler code coverage** across 24-hour fuzzing trials
 - Computed basic block coverage via the AFL-QEMU-Cov utility

Metrics	Bin2Wrong vs. Cornucopia	Bin2Wrong vs. DecFuzzer
Basic Block Coverage	+16%	+32%
Generated Binaries	-74%	-95%

- Takeaways:** Bin2Wrong's binaries each **exercise more decompiler internals**

Evaluation: Decompiler Bugs

- Enumerated **total unique bugs** throughout all 24-hour fuzzing campaigns
 - All bugs manually deduplicated and reported to their respective decompiler developers

Metrics	DecFuzzer	Cornucopia
Total Found Bugs	0	10
Individually-found	0	4
Confirmed or fixed	0	5

Decompiler	DecFuzzer	Cornucopia
Angr	0	2
BinNinja	0	0
Reko	0	3
R2Ghidra	0	1
Relyze	0	2
RetDec	0	2
rev.ng	0	0

Evaluation: Decompiler Bugs




- Enumerated **total unique bugs** throughout all 24-hour fuzzing campaigns
 - All bugs manually deduplicated and reported to their respective decompiler developers

Metrics	DecFuzzer	Cornucopia	Bin2Wrong
Total Found Bugs	0	10	48
Individually-found	0	4	42
Confirmed or fixed	0	5	30

Decompiler	DecFuzzer	Cornucopia	Bin2Wrong
Angr	0	2	9
BinNinja	0	0	11
Reko	0	3	6
R2Ghidra	0	1	2
Relyze	0	2	7
RetDec	0	2	11
rev.ng	0	0	2

- Takeaways:** Bin2Wrong's diverse binaries **expose more decompiler bugs**
 - Since reporting, **30 are now confirmed and/or fixed**

Case Studies: Bin2Wrong-found Bugs



Case Studies: Bin2Wrong-found Bugs

- Bin2Wrong found a **critical bug** in commercial decompiler **Binary Ninja**

```
int var = 0, int idx = 1;
switch (var) {
    case 0: var=5; break;
    case 1: ...
    case 2: ...
    default: idx=0; break;
}
```

Results: var=?, idx=?




```
int var = 0, int idx = 1;
if (var == 2) { ... }
else {
    if (var == 0): var=5;
    if (var == 1): ...
    if (var > 2): idx=0;
}
```

Results: var=?, idx=?

Case Studies: Bin2Wrong-found Bugs

- Bin2Wrong found a **critical bug** in commercial decompiler **Binary Ninja**



Case Studies: Bin2Wrong-found Bugs

- Bin2Wrong found a **critical bug** in commercial decompiler **Binary Ninja**
 - Spurred **an overhaul** of Binary Ninja's internal control flow recovery processes

The image shows a composite view. On the left, a GitHub issue page for 'binaryninja-api' is displayed, showing a closed bug related to HLIL control flow structuring. On the right, a blog post titled 'RESTRUCTURING THE BINARY NINJA DECOMPILER' is shown, dated June 19, 2024, by Rusty Wagner. The blog post discusses a privately reported flaw and expresses gratitude to researchers.

binary.ninja/2024/06/19/restructuring-the-decompiler.html

- Takeaways:** Bin2Wrong **exposes critical decompiler bugs missed by others**

Conclusion: Why Bin2Wrong?

- **Decompiler errors** create **downstream failures**
 - Hence, **testing decompilers** is critical to fixing them
- Prior fuzzers **fail to thoroughly test** decompilers
 - Only partial coverage of source/compilation factors
 - Vast majority of binary diversity thus left unexplored

Conclusion: Why Bin2Wrong?

- **Decompiler errors** create **downstream failures**
 - Hence, **testing decompilers** is critical to fixing them
- Prior fuzzers **fail to thoroughly test** decompilers
 - Only partial coverage of source/compilation factors
 - Vast majority of binary diversity thus left unexplored
- Our solution: **Bin2Wrong**
 - Unified mutation **coalescing source and compilation**
 - Support for **6 compilers**, **5,183 optimizations**, all major **executable formats**, and **virtually all C code constructs**
 - **Outcome:** systematic decompiler fuzzer exploring all of these individual dimensions—and **combinations thereof**

Key Results:

6–16× higher binary diversity,
16–32% higher code coverage,
48 new errors,
30 confirmed

Thank you!



github.com/FuturesLab/Bin2Wrong

Contact:

✉ zao.yang@utah.edu

𝕏 @yangzaocn

𝕏 @zaoyang.bsky.social

FUTURES³
LAB FUTURE TECHNOLOGY FOR USABLE, RELIABLE, &
EFFICIENT SECURITY OF SOFTWARE & SYSTEMS