

No Harness, No Problem:

Oracle-guided Harnessing for Auto-generating C API Fuzzing Harnesses

Gabriel Sherman

University of Utah

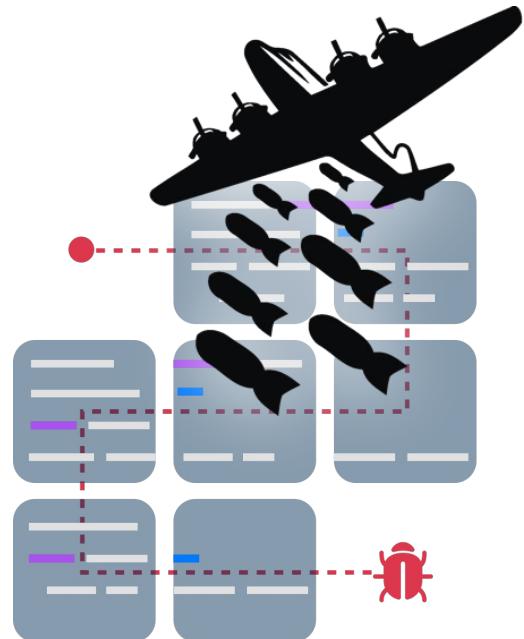
Stefan Nagy

University of Utah



Fuzzing: Automated Bug Discovery

- Today's leading **automated bug-finding strategy**
- Fundamental steps:
 - Generate many inputs via randomized **mutation**
 - Execute** and check results (**crashes**, **hangs**, etc.)
 - New coverage? **Save and re-mutate** that input!



Fuzzing: Automated Bug Discovery

- Today's leading **automated bug-finding strategy**
- Fundamental steps:
 - Generate many inputs via randomized **mutation**
 - Execute** and check results (**crashes**, **hangs**, etc.)
 - New coverage? **Save and re-mutate** that input!
- Revealed **thousands of bugs** in modern software
 - Apps and libraries
 - Operating systems
 - IoT devices and more!



Harnesses: the Foundation of Fuzzing

- Wrapper programs that feed fuzzer input to API functions

```
int harness(const uint8_t *Data, size_t Size){  
    initGEOS(geos_msg_handler, geos_msg_handler);  
    Reader *reader = Reader_create();  
    Geometry *g1 = Reader_read(reader, Data);  
    if (g1 != NULL) {  
        Geometry *g2 = FromWKB_buf(Data, Size);  
        if (g2 != NULL) {  
            Geometry *g3 = Intersection(g1, g2);  
            Geom_destroy(g3);  
            Geom_destroy(g2);  
        }  
        ...  
    }  
}
```

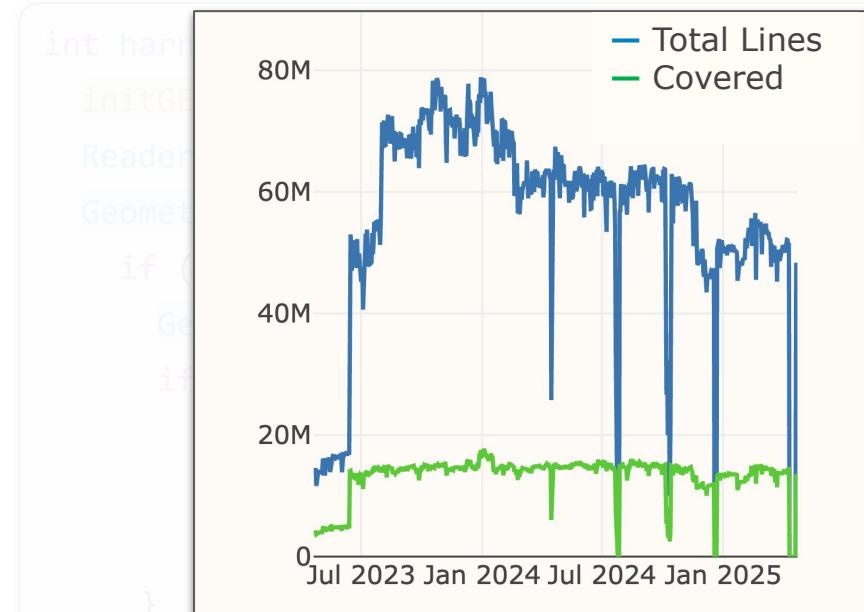
Harnesses: the Foundation of Fuzzing

- Wrapper programs that feed fuzzer input to API functions
- Each part serves a purpose:
 - Library initialization
 - Data entrypoints
 - Data processing

```
int harness(const uint8_t *Data, size_t Size){  
    initGEOS(geos_msg_handler, geos_msg_handler);  
    Reader *reader = Reader_create();  
    Geometry *g1 = Reader_read(reader, Data);  
    if (g1 != NULL) {  
        Geometry *g2 = FromWKB_buf(Data, Size);  
        if (g2 != NULL) {  
            Geometry *g3 = Intersection(g1, g2);  
            Geom_destroy(g3);  
            Geom_destroy(g2);  
        }  
        ...  
    }  
}
```

Harnesses: the Foundation of Fuzzing

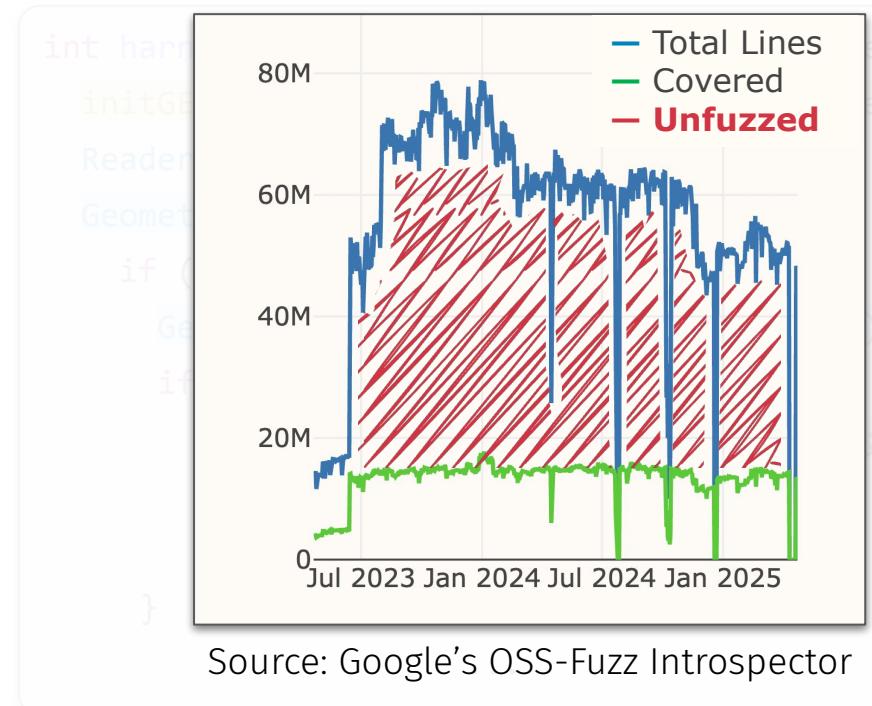
- Wrapper programs that feed fuzzer input to API functions
- Each part serves a purpose:
 - Library initialization
 - Data entrypoints
 - Data processing
- Problem:** so much of today's code remains **un-harnessed**



Source: Google's OSS-Fuzz Introspector

Harnesses: the Foundation of Fuzzing

- Wrapper programs that feed fuzzer input to API functions
- Each part serves a purpose:
 - Library initialization
 - Data entrypoints
 - Data processing
- Problem:** so much of today's code remains **un-harnessed**
 - Hence, it is **not being fuzzed**



Source: Google's OSS-Fuzz Introspector

Prior Work: Automatic Harness Generation

- **Idea:** create harnesses automatically exercising diverse API usage

How API Usage is Learned:	Advantages:	Disadvantages:
Available API reference code (Fudge, FuzzGen, Utopia, Rubick)	Generate harnesses fast	Reference code covers very little

Prior Work: Automatic Harness Generation

- **Idea:** create harnesses automatically exercising diverse API usage

How API Usage is Learned:	Advantages:	Disadvantages:
Available API reference code (Fudge, FuzzGen, Utopia, Rubick)	Generate harnesses fast	Reference code covers very little
Pre-written API specification (GraphFuzz)	Higher semantic validity	Specification writing is unscalable

Prior Work: Automatic Harness Generation

- Idea: create harnesses automatically exercising diverse API usage

How API Usage is Learned:	Advantages:	Disadvantages:
Available API reference code (Fudge, FuzzGen, Utopia, Rubick)	Generate harnesses fast	Reference code covers very little
Pre-written API specification (GraphFuzz)	Higher semantic validity	Specification writing is unscalable
On-the-fly random mutation (Hopper)	Harnesses far more code	Broken semantics = false crashes

Prior Work: Automatic Harness Generation

- Idea: create harnesses automatically exercising diverse API usage
 - Though **Hopper** eschews specifications or reference code, it often breaks API semantics

```
void handler(const char *fmt, ...){  
    exit(EXIT_FAILURE);  
}  
int main(){  
    ContextHandle_HS *c = GEOS_init_r();  
    Context_setErrorHandler_r(c, handler);  
    OrientPolygons_r(c, ...); // Crash!  
}
```

Missed Initialization

```
ContextHandle_HS *c = GEOS_init_r();  
Geom_t *x = Geom_createEmptyPoint_r(c);  
Geom_t *y = Geom_createCollection_r(c,-9,&x,6);  
if (y == NULL) return 0;  
Normalize_r(c, y); // Crash!
```

Missed Sanity Checks

```
ContextHandle_HS *c = GEOS_init_r();  
WKBReader_t *r = WKBReader_create_r(c);  
WKBReader_destroy(r);  
WKBReader_destroy_r(c, r); // Crash!
```

Erroneous Call Sequences

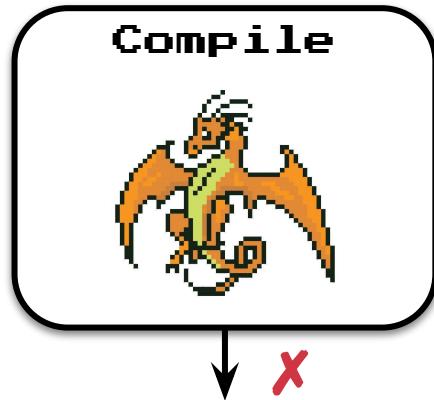
On-the-fly randomization
(Hopper)

= false crashes

Motivation: how can we harness **correctly**—without turning back to **specifications** or **reference code**?

Inspiration: Harnessing by Hand

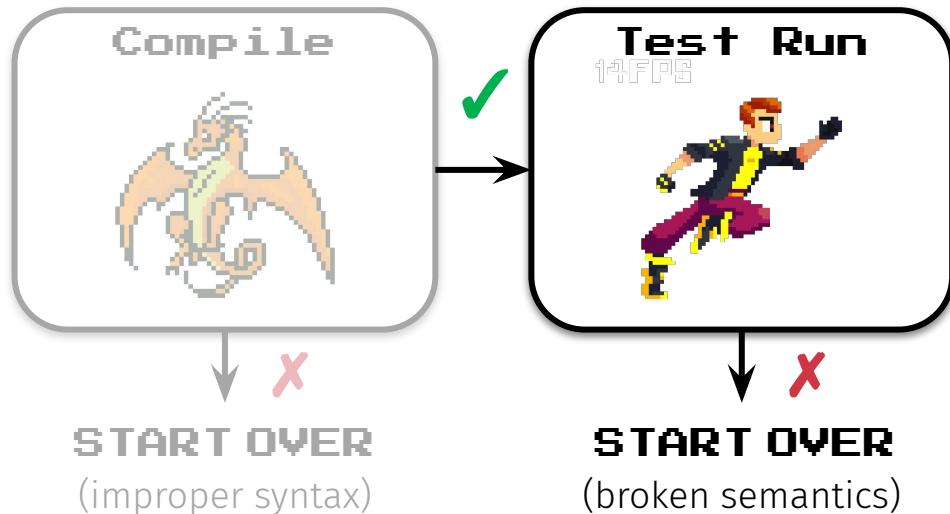
- Manual harness development spans **three sequential steps**
 - Failures at any step signal the harness is **wrong**, prompting **refinement**



START OVER
(improper syntax)

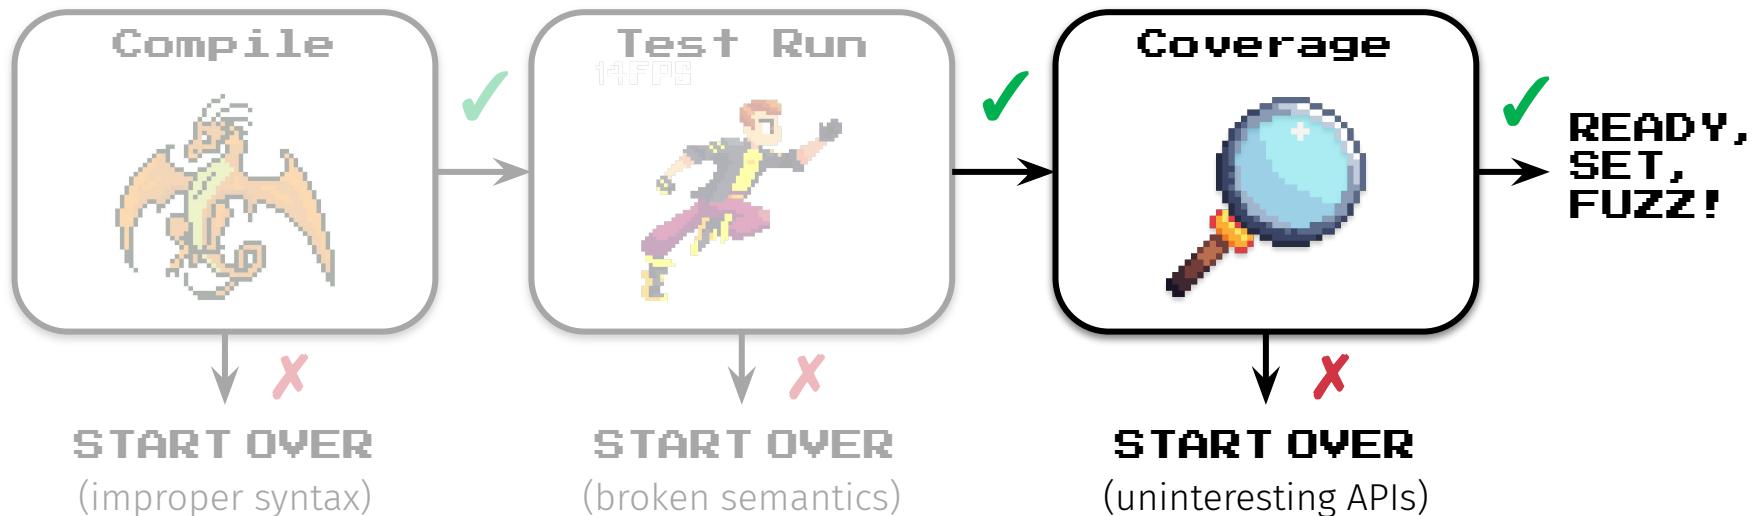
Inspiration: Harnessing by Hand

- Manual harness development spans **three sequential steps**
 - Failures at any step signal the harness is **wrong**, prompting **refinement**



Inspiration: Harnessing by Hand

- Manual harness development spans **three sequential steps**
 - Failures at any step signal the harness is **wrong**, prompting **refinement**



Inspiration: Harnessing by Hand

- Manual harness development spans **three sequential steps**
 - Failures at any step signal the harness is **wrong**, prompting **refinement**

Observation: **compilation**, **execution**, and **coverage** serve as **oracles** in manual harness development...

Our goal: leverage these to harness **automatically!**

START OVER
(improper syntax)

START OVER
(broken semantics)

START OVER
(uninteresting APIs)

READY,
SET,
FUZZ!

Our Solution: Oracle-guided Harnessing

- **Idea:** mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**

Our Solution: Oracle-guided Harnessing

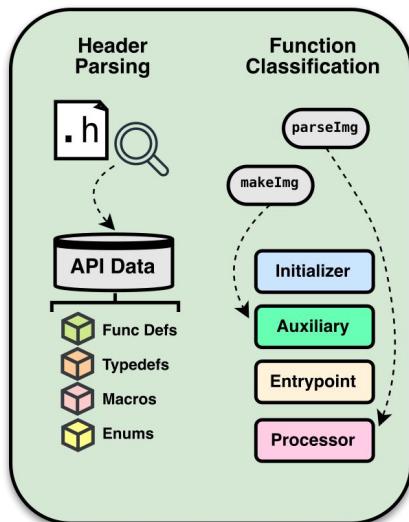
- **Idea:** mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**
- **Insight:** all functions serve a purpose:
 - Surveyed OSS-Fuzz's **281** C libraries:
 - Library **initialization** routines
 - Data **entrypoints**, **processors**
 - Data-resolving **auxiliary calls**
- **Approach:** build harnesses that follow these common-case API patterns!

```
int harness(const uint8_t *Data, size_t Size){  
    initGEOS(geos_msg_handler, geos_msg_handler);  
    Reader *reader = Reader_create();  
    Geometry *g1 = Reader_read(reader, Data);  
    if (g1 != NULL) {  
        Geometry *g2 = FromWKB_buf(Data, Size);  
        if (g2 != NULL) {  
            Geometry *g3 = Intersection(g1, g2);  
            ...  
        }  
    }  
}
```

Our Solution: Oracle-guided Harnessing

- **Idea:** mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**

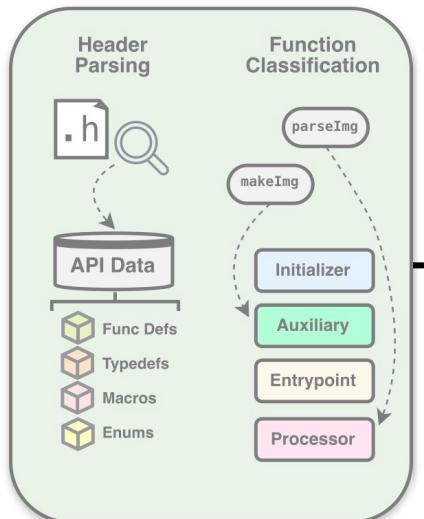
1. Pre-process Headers



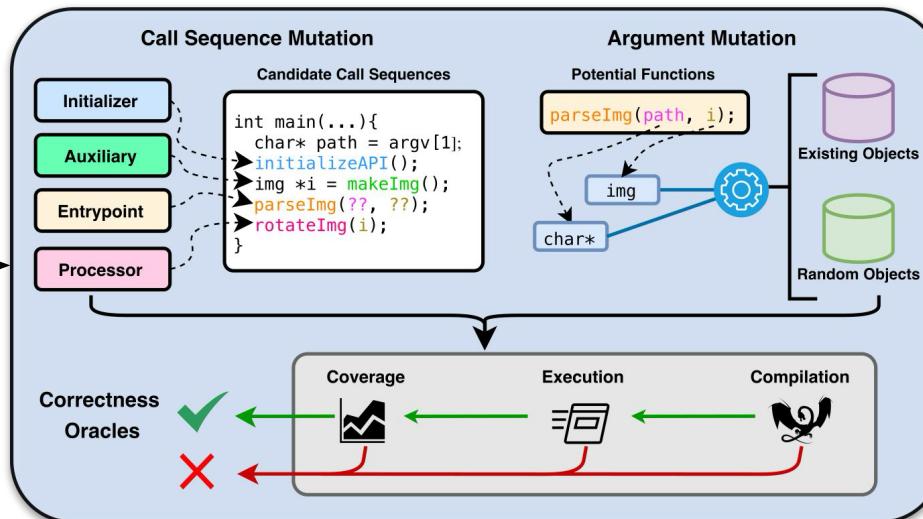
Our Solution: Oracle-guided Harnessing

- Idea: mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**

1. Pre-process Headers



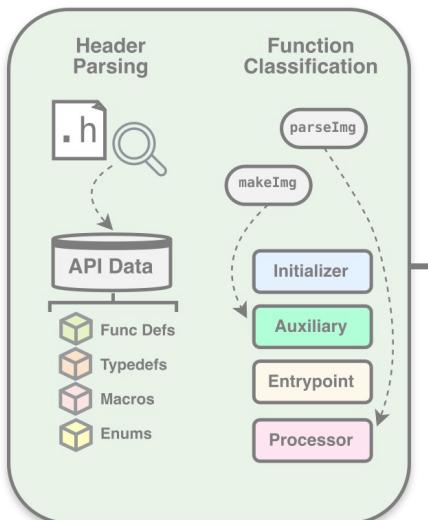
2. Iterative Harness Synthesis



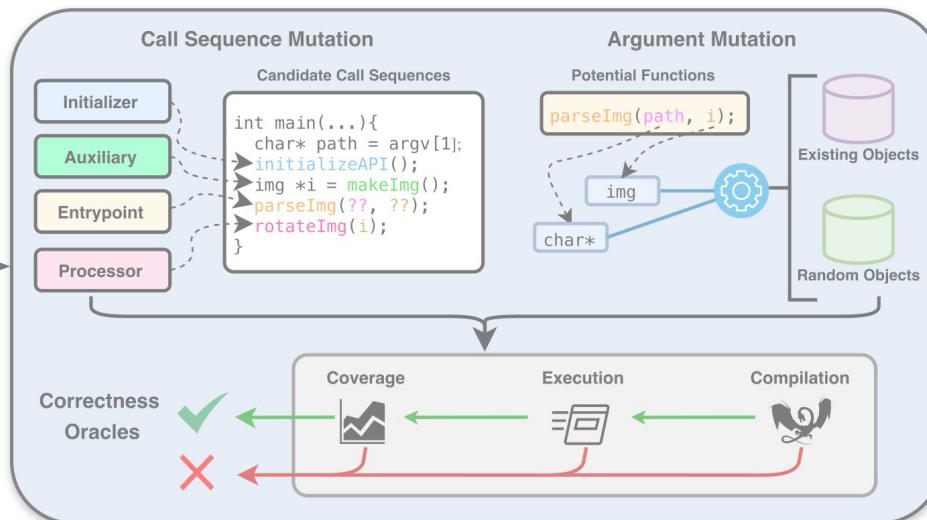
Our Solution: Oracle-guided Harnessing

- Idea: mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**

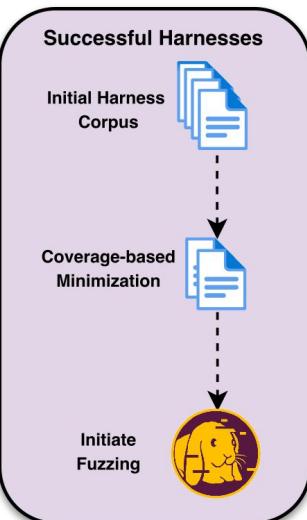
1. Pre-process Headers



2. Iterative Harness Synthesis

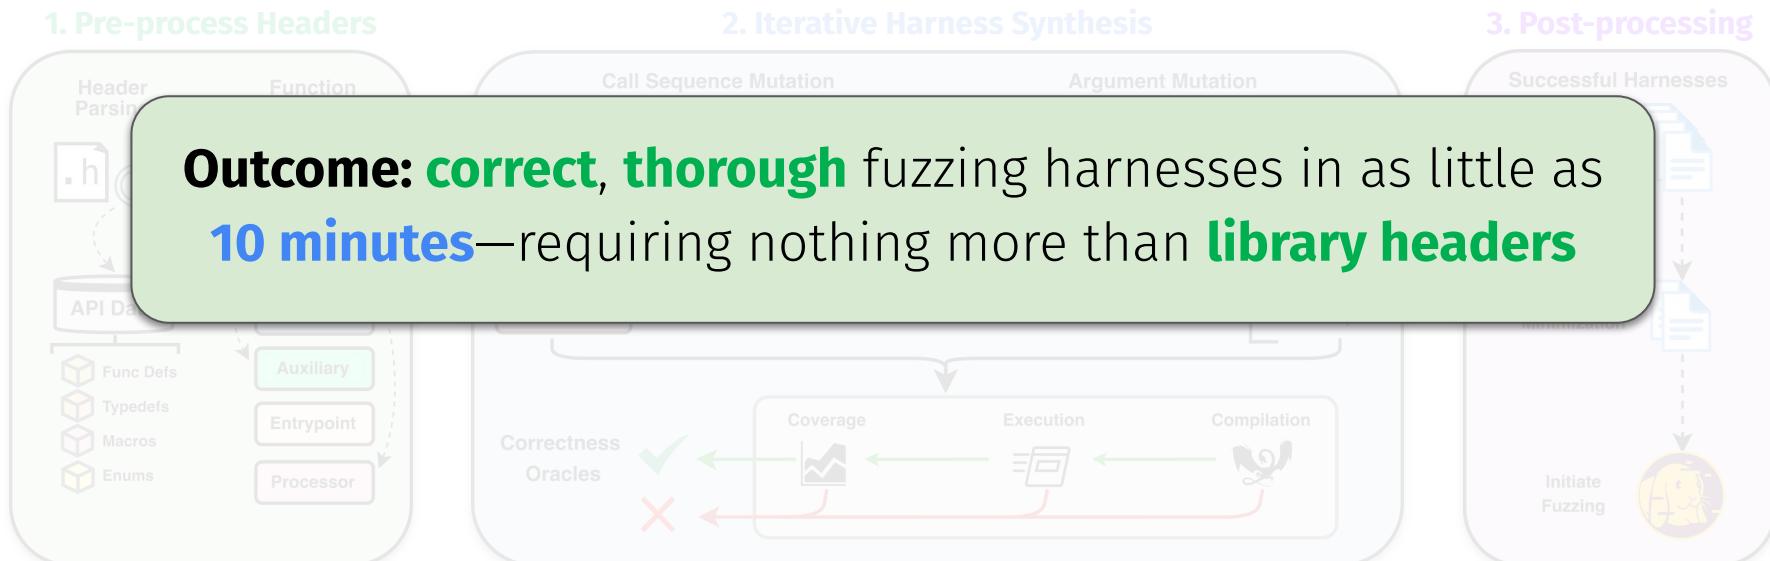


3. Post-processing



Our Solution: Oracle-guided Harnessing

- Idea: mutationally build-up harnesses, validating each step via our **oracles**
 - Culling **invalid** or **uninteresting** API usage gradually converges on **valid harnesses**



Evaluation: Overview

- Prototype: **OGHarn**
 - Supports harnessing for C libraries
- Evaluated on **20 real-world APIs**
 - **16** with prior harnesses, **4** without
 - Coverage, correctness bug-finding
 - All campaigns run for 5x24-hr trials
- Competitors:
 - Hopper (current state-of-the-art)
 - Existing OSS-Fuzz harnesses (**33** total)

Libraries with existing harnesses:	C-Ares, cglTF, cJSON, GPAC, HDF5, LCMS, Lexbor, libGEOS, libICAL, libMagic, libPCAP, libUCL, OpenEXR, PCRE2, SQLite3, Zlib
Libraries without prior harnesses:	Faup, libFYAML, RayLib, StormLib

Evaluation: Correctness

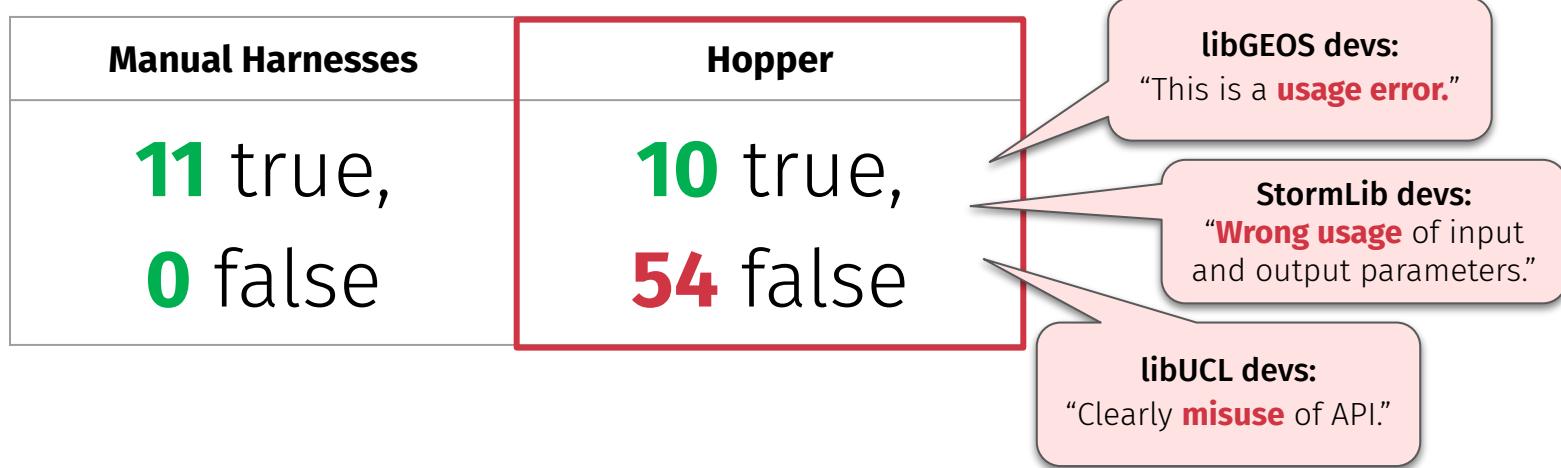
- Measured **false-positive bugs** found by all competitors:
 - Initial evaluation on OGHarn's **top-10 highest-coverage** harnesses

Manual Harnesses

11 true,
0 false

Evaluation: Correctness

- Measured **false-positive bugs** found by all competitors:
 - Initial evaluation on OGHarn's **top-10 highest-coverage** harnesses



Evaluation: Correctness

- Measured **false-positive bugs** found by all competitors:
 - Initial evaluation on OGHarn's **top-10 highest-coverage** harnesses
 - Scaled-up to include **all harnesses**; resulted in one **new bug found**

Manual Harnesses	Hopper	Oracle-guided Harnessing
11 true, 0 false	10 true, 54 false	41 true, 0 false

- Takeaways:** Oracle-guided Harnessing correctly **upholds API semantics**

Evaluation: Code Coverage

- Measured OGHarn's relative **total and unique code coverage**
 - Median **653% more functions harnessed** over developer-written harnesses

Versus Manual Harnesses
+14% total,
+364% unique

Evaluation: Code Coverage

- Measured OGHarn's relative **total and unique code coverage**
 - Median **653% more functions harnessed** over developer-written harnesses

Versus Manual Harnesses	Versus Hopper
+14% total, +364% unique	-31% total, -72% unique

- Takeaways:** Hopper covers more, but at steep expense of **correctness**

Evaluation: New Bugs Found

- Quantified all competitors' **total and uniquely-found bugs**
 - Following our own triage, we reported newly-found bugs to developers

Manual Harnesses	Hopper
11 total, 4 unique	10 total, 8 unique

Evaluation: New Bugs Found

- Quantified all competitors' **total and uniquely-found bugs**
 - Following our own triage, we reported newly-found bugs to developers

Manual Harnesses	Hopper	Oracle-guided Harnessing
11 total, 4 unique	10 total, 8 unique	41 total, 32 unique

- Takeaways:** Oracle-guided Harnessing **expands fuzzing bug discovery**
 - All **41 are confirmed**, with **40 patched** post-reporting

Conclusion: Why Oracle-guided Harnessing?

- Prior harnessing approaches **restricted** by need for **reference code** or **specifications**
- Recent workarounds fail to preserve **valid API usage**, causing **false-positive crashes**

Conclusion: Why Oracle-guided Harnessing?

- Prior harnessing approaches **restricted** by need for **reference code or specifications**
- Recent workarounds fail to preserve **valid API usage**, causing **false-positive crashes**
- Our solution: **Oracle-guided Harnessing**
 - Mutational harness creation **using only headers**
 - Structure generation by **common-case patterns**
 - Leverage oracles—**compilation, execution, and code coverage**—to validate harness mutations
 - **Outcome:** valid, thorough API fuzzing harnesses

Key Results:

41 new bugs,

0 false bugs,

14% coverage
over OSS-Fuzz

Thank you!



Try Oracle-guided Harnessing at:
github.com/FuturesLab/OGHarn

Contact: gabe.sherman@utah.edu
gabe-sherman.github.io
futures.cs.utah.edu

FUTURES³
LAB FUTURE TECHNOLOGY FOR USABLE, RELIABLE, &
EFFICIENT SECURITY OF SOFTWARE & SYSTEMS